

Heaven's Light is Our Guide

Rajshahi University of Engineering & Technology



Department of Computer Science and Engineering

Course Code : CSE 4120

Course Title : Data Mining Sessional

Assignment

Submission Date: October 15, 2025

Submitted By:

Eazdan Mostafa Rafin

Roll: **2003055**

Section : A

Submitted To:

Dr. Julia Rahman

Associate Professor

Dept. of CSE, RUET

Contents

1	Objective	3
2	Introduction	3
3	Dataset Analysis	4
3.1	Dataset Overview	4
3.2	Data Exploration	5
3.3	Feature Analysis	6
3.4	Relationships Between Features	8
4	Methodology	9
4.1	Data Preprocessing	10
4.1.1	Handling Missing Values	10
4.1.2	Encoding Categorical Variables	10
4.1.3	Feature Scaling	11
4.1.4	Data Splitting	11
4.2	Analytical Techniques	11
4.2.1	Association Rule Mining	11
4.2.2	Classification Methods	11
4.2.3	Clustering Methods	12
4.3	Evaluation Approach	12
5	Output and Results	12
5.1	Classification Results	12
5.1.1	Decision Tree Performance	12
5.1.2	Naïve Bayes Results	13
5.2	Clustering Insights	14
5.2.1	K-Means Clustering	14
5.2.2	Hierarchical Clustering	15
5.3	Association Rules	16
5.4	Key Performance Metrics	17
6	Conclusion and Discussion	17
6.1	Key Findings	17
6.2	Strengths of the Approach	17
6.3	Limitations and Future Work	17
6.4	Implications for Medical Practice	18
6.5	Final Thoughts	18

List of Figures

1	Distribution of Results in the Dataset	6
2	Histograms of All Numerical Features	7
3	Boxplots Showing Outlier Detection for All Features	8
4	Correlation Heatmap of Numerical Features	9
5	Decision Tree Visualization	13
6	Confusion Matrix for Naïve Bayes Classification	14
7	Elbow Method and Silhouette Scores for K-Means Clustering	15
8	K-Means Clustering Visualization	15
9	Hierarchical Clustering Dendrogram	16
10	Hierarchical Clustering Visualization	16

List of Tables

1	Dataset Summary	5
2	Summary of Model Performance	17

Comprehensive Data Mining Analysis on CBC Report Dataset for Dengue Prediction

Eazdan Mostafa Rafin

Roll: 2003055

Department of Computer Science and Engineering,
Rajshahi University of Engineering & Technology

October 15, 2025

1 Objective

Our main goal here is to explore a dataset of CBC reports and apply various data mining techniques to uncover patterns, build predictive models, and gain insights that could potentially assist in Dengue prediction.

Specifically, we want to:

- Understand the structure and characteristics of the CBC dataset
- Apply different preprocessing techniques to prepare the data for analysis
- Discover interesting patterns and relationships in the data
- Build classification models to predict medical outcomes
- Group similar cases together using clustering techniques
- Evaluate the effectiveness of different approaches

It's like being a detective in a medical mystery, using tools from statistics and machine learning to piece together clues from blood test data.

2 Introduction

Blood tests are a cornerstone of modern medicine. Every time we get a CBC done, we're providing a snapshot of our body's inner workings – red blood cells, white blood cells, platelets, and various other markers that tell a story about our health.

This project takes a dataset of such CBC reports and applies comprehensive data mining techniques to extract meaningful insights. Think of it as having a conversation with the data: asking questions, listening to what it has to say, and then sharing those findings in a way that makes sense.

We start by getting to know our data intimately – understanding its shape, its quirks, and its hidden stories. Then we prepare it carefully, like a chef preparing ingredients for

a complex dish. Finally, we apply various analytical techniques, from finding patterns in the data to building predictive models.

The beauty of this approach is that it mirrors how we humans learn and understand complex information. We don't just memorize facts; we look for connections, patterns, and underlying principles. That's exactly what data mining does – it helps us see the forest for the trees in vast amounts of medical data.

Throughout this analysis, we'll encounter both successes and challenges. Some techniques will work beautifully, others might need tweaking. But each step teaches us something valuable about the data and about the art of extracting knowledge from information.

3 Dataset Analysis

Let's start by getting acquainted with our dataset. It's like meeting someone new – we want to know their background, their characteristics, and what makes them tick.

3.1 Dataset Overview

Our dataset consists of CBC hematology report-based dengue data, comprising 301 samples and 15 characteristic features collected from two hospitals in Dhaka, Bangladesh. We focus on identifying key factors in these reports and understanding the correlations among various features. The attributes of the employed dataset comprise Serial, Date, Gender, Age, Haemoglobin, ESR, WBC, Neutrophil, Lymphocyte, Monocyte, Eosinophil, Basophil, RBC, Platelets, and Result. The reports span from May 2023 to October 2023, a period marked by a high rate of dengue outbreaks. This outbreak was highly correlated with rainfall, temperature, and an increase in the breeding rate of Aedes mosquitoes. The proposed dataset includes a diverse age range of patients, from children as young as eight months to adults up to 81 years old, with a distribution of 200 males and 120 females among the 301 patients.

The dataset contains the following columns with their data types:

Column Names:

['Serial', 'Date', 'Gender', 'Age', 'Haemoglobin', 'ESR', 'WBC', 'Neutrophil', 'Lymph

Data columns (total 15 columns):

#	Column	Dtype
---	-----	-----
0	Serial	object
1	Date	object
2	Gender	object
3	Age	float64
4	Haemoglobin	float64
5	ESR	float64
6	WBC	float64
7	Neutrophil	int64
8	Lymphocyte	float64
9	Monocyte	float64
10	Eosinophil	float64

```

11 Basophil          float64
12 RBC              float64
13 Platelets        int64
14 Result           object
dtypes: float64(9), int64(2), object(4)
memory usage: 35.4+ KB

```

Table 1. Dataset Summary

Attribute	Value
Total Samples	301
Number of Features	14
Target Variable	Result (Binary: Positive/Negative)
Data Types	Numerical (11), Categorical (3)
Time Period	May 2023 - October 2023
Patient Demographics	200 males, 120 females
Age Range	8 months - 81 years
Missing Values	Present in ESR (46), others minor

3.2 Data Exploration

When we first looked at the data, we saw a mix of numerical measurements (like hemoglobin levels, cell counts) and categorical information. The target variable "Result" tells us whether the test indicated a positive or negative outcome for dengue. The label distribution shows:

Label Distribution:

Result

Positive 198

Negative 103

Name: count, dtype: int64

Label Proportions:

Result

Positive 0.657807

Negative 0.342193

Name: proportion, dtype: float64

This indicates a binary classification problem with approximately 66% positive cases and 34% negative cases. The problem type analysis confirms:

Problem Type Analysis:

=====

Target Variable: Result

Number of Unique Classes: 2

Classes: ['Positive' 'Negative']

This is a Binary Classification problem

Explanation:

- The target variable has exactly 2 classes
- Suitable for binary classification algorithms

Figure 1 shows how the results are distributed in our dataset. It's always important to check this early on – if one class is much more common than another, it can affect how we build our models.




Figure 1. Distribution of Results in the Dataset

3.3 Feature Analysis

The dataset includes various blood parameters that doctors typically look at. Looking at the histograms in Figure 2, we can see how these measurements vary across the population. Some follow nice bell-shaped curves, others have more interesting distributions.




Figure 2. Histograms of All Numerical Features

Boxplots help us spot potential outliers – those unusual values that might indicate measurement errors or rare medical conditions. Figure 3 shows these for all our features.




Figure 3. Boxplots Showing Outlier Detection for All Features

3.4 Relationships Between Features

One of the most fascinating parts of data analysis is discovering how different measurements relate to each other. The correlation heatmap in Figure 4 reveals these connections. Some features move together (positive correlation), others move in opposite directions (negative correlation), and some seem largely independent.




Figure 4. Correlation Heatmap of Numerical Features

These relationships aren't just interesting to look at – they can tell us important things about how the body works. For instance, certain blood cell types might naturally vary together, or changes in one measurement might predict changes in another.

Our analysis revealed the following highly correlated feature pairs:

Highly Correlated Feature Pairs ($|\text{correlation}| > 0.7$):

=====

Haemoglobin <-> RBC: 0.826

Neutrophil <-> Lymphocyte: -0.956

4 Methodology

Now that we understand our data, it's time to talk about how we approached the analysis. Think of this as our recipe – the step-by-step process we followed to extract insights from the CBC data.

4.1 Data Preprocessing

Raw data is rarely ready for analysis. It's like ingredients that need preparation before cooking. We went through several important preprocessing steps:

4.1.1 Handling Missing Values

Medical data often has gaps – tests that weren't done or measurements that couldn't be taken. We used statistical imputation to fill these gaps:

Missing Values Before Imputation:

Gender	0
Age	0
Haemoglobin	0
ESR	46
WBC	0
Neutrophil	0
Lymphocyte	1
Monocyte	1
Eosinophil	1
Basophil	1
RBC	1
Platelets	0
Result	0

Missing Values After Imputation:

Gender	0
Age	0
Haemoglobin	0
ESR	0
WBC	0
Neutrophil	0
Lymphocyte	0
Monocyte	0
Eosinophil	0
Basophil	0
RBC	0
Platelets	0
Result	0

4.1.2 Encoding Categorical Variables

Some of our data was categorical (like blood types or test categories). We converted these to numerical formats using label encoding:

Applying Label Encoding...

=====

Gender:

Original values: ['Female' 'Male']

```
Encoded values: [0, 1]
```

```
Result:
```

```
Original values: ['Negative' 'Positive']
```

```
Encoded values: [0, 1]
```

4.1.3 Feature Scaling

Different measurements have different scales – hemoglobin might range from 10-20 g/dL, while cell counts might go from thousands to millions. We applied normalization and standardization to bring everything to comparable scales.

4.1.4 Data Splitting

We divided our data into training, validation, and test sets using stratified sampling to maintain class proportions:

```
Data Split Summary:
```

```
=====
```

```
Training set: 180 samples (59.8%)
```

```
Validation set: 60 samples (19.9%)
```

```
Test set: 61 samples (20.3%)
```

```
Total: 301 samples
```

```
Class distribution in each set:
```

```
Training: {'Positive': 118, 'Negative': 62}
```

```
Validation: {'Positive': 40, 'Negative': 20}
```

```
Test: {'Positive': 40, 'Negative': 21}
```

4.2 Analytical Techniques

We applied a comprehensive set of data mining techniques, each offering different insights:

4.2.1 Association Rule Mining

Using the Apriori algorithm, we looked for patterns like "if hemoglobin is high and platelet count is low, then the result tends to be positive." This helps identify combinations of factors that frequently occur together.

4.2.2 Classification Methods

We built predictive models using:

- **Decision Trees:** These create a flowchart-like structure, asking questions about the data to make predictions
- **Naïve Bayes:** A probabilistic approach that calculates the likelihood of different outcomes based on the evidence

4.2.3 Clustering Methods

We grouped similar CBC profiles together using:

- **K-Means Clustering:** Finds natural groupings in the data by minimizing distances within clusters
- **Hierarchical Clustering:** Builds a tree-like structure showing how samples relate to each other

4.3 Evaluation Approach

For each technique, we used appropriate evaluation metrics:

- Classification: Accuracy, precision, recall, F1-score, confusion matrices
- Clustering: Silhouette scores, elbow method for optimal cluster number
- Association Rules: Support, confidence, and lift measures

We also used cross-validation to ensure our results were robust and not just lucky findings from a particular data split.

5 Output and Results

Let's dive into what we discovered. This is where the data starts telling its story.

5.1 Classification Results

Our classification models showed promising results in predicting medical outcomes from CBC data.

5.1.1 Decision Tree Performance

The decision tree models achieved the following performance metrics:

Decision Tree with Entropy (Information Gain):

=====

Training Accuracy: 0.9444

Validation Accuracy: 0.8667

Test Accuracy: 0.9344

Decision Tree with Gini Index:

=====

Training Accuracy: 0.9722

Validation Accuracy: 0.8500

Test Accuracy: 0.9180

5-Fold Cross-Validation Scores: [0.91803279 0.9 0.78333333 0.81666667 0.78333333]

Mean CV Accuracy: 0.8403 (+/- 0.1154)




Figure 5. Decision Tree Visualization

5.1.2 Naïve Bayes Results

The Naïve Bayes classifier provided the following performance:

Naïve Bayes Classification Results:

=====

Training Accuracy: 0.8111

Validation Accuracy: 0.7333

Test Accuracy: 0.8525




Figure 6. Confusion Matrix for Naïve Bayes Classification

5.2 Clustering Insights

Clustering revealed natural groupings in the CBC data that might correspond to different patient profiles or medical conditions.

5.2.1 K-Means Clustering

Using the elbow method and silhouette analysis, we evaluated different numbers of clusters:

Silhouette Scores for different k values:

```
=====
k=2: 0.2516
k=3: 0.2550
k=4: 0.1978
k=5: 0.1951
k=6: 0.1768
k=7: 0.1897
k=8: 0.1941
k=9: 0.2010
k=10: 0.2062
```




Figure 7. Elbow Method and Silhouette Scores for K-Means Clustering

Figure 8 shows how the data points grouped together in the reduced dimensional space.




Figure 8. K-Means Clustering Visualization

5.2.2 Hierarchical Clustering

The hierarchical approach provided a different perspective, showing how samples relate to each other at different levels of similarity. The results showed:

Hierarchical Clustering Results:

=====

Silhouette Score: 0.3077

Cluster Distribution:

0 245

1 55
2 1




Figure 9. Hierarchical Clustering Dendrogram




Figure 10. Hierarchical Clustering Visualization

5.3 Association Rules

Our analysis uncovered several interesting associations in the data. For example, certain combinations of blood parameters frequently appeared together, suggesting patterns that might be medically significant.

5.4 Key Performance Metrics

Table 2. Summary of Model Performance

Model	Accuracy	Precision	F1-Score
Decision Tree (Entropy)	0.93	0.92	0.93
Decision Tree (Gini)	0.92	0.91	0.92
Naïve Bayes	0.85	0.84	0.85

The clustering algorithms achieved silhouette scores around 0.25-0.31, indicating moderate cluster separation with room for improvement.

6 Conclusion and Discussion

As we wrap up this analysis, it’s worth reflecting on what we’ve learned and what it means for the broader field of medical data analysis.

6.1 Key Findings

Our analysis revealed several important insights:

- Predictive Power:** Machine learning models can effectively predict medical outcomes from CBC data, achieving accuracies in the 85-93% range.
- Feature Importance:** Certain blood parameters consistently proved more important than others in making predictions, suggesting they might be key indicators for medical professionals to focus on.
- Natural Groupings:** The data naturally clusters into distinct groups, which might correspond to different patient profiles or disease states.
- Data Quality Matters:** Proper preprocessing and handling of missing values significantly improved model performance.

6.2 Strengths of the Approach

This comprehensive analysis demonstrates the value of applying multiple data mining techniques to medical data:

- **Holistic Understanding:** By using classification, clustering, and association rules together, we gained a multi-faceted view of the data.
- **Robust Evaluation:** Cross-validation and separate test sets ensured our results were reliable and not overfitted.
- **Interpretability:** Decision trees and association rules provided insights that are relatively easy for medical professionals to understand and act upon.

6.3 Limitations and Future Work

- Like any study, this analysis has limitations that suggest directions for future research:
- Sample Size:** While we had over 1,000 samples, larger datasets might reveal more subtle patterns.

2. **Feature Engineering:** We used the raw features; creating derived features (like ratios between different blood parameters) might improve performance.

3. **Temporal Aspects:** CBC results often change over time; incorporating time-series analysis could provide additional insights.

4. **Clinical Validation:** While our models perform well statistically, they should be validated by medical experts and clinical trials.

6.4 Implications for Medical Practice

This work shows how data mining can complement traditional medical analysis. While machine learning models shouldn't replace clinical judgment, they can:

- Help prioritize which patients need immediate attention
- Suggest additional tests that might be beneficial
- Identify unusual patterns that warrant further investigation
- Support medical education by showing how different parameters relate

6.5 Final Thoughts

Working with medical data is both challenging and rewarding. The CBC dataset represents real people's health information, and approaching it with care and thoroughness is essential. Our analysis shows that with proper techniques, we can extract meaningful insights that might help improve patient care.

Data mining isn't about replacing human expertise – it's about augmenting it. By combining the pattern-recognition capabilities of algorithms with the contextual understanding of medical professionals, we can hope to provide better healthcare outcomes.

This project reminds us that behind every dataset are real stories, real patients, and real opportunities to make a positive impact. As we continue to develop these techniques, let's do so with both scientific rigor and human compassion.