

Heaven's Light is Our Guide

Rajshahi University of Engineering & Technology



Department of Computer Science and Engineering

Course Code : CSE 4120

Course Title : Data Mining Sessional

Lab Report - 2

Experiment Name: Data Preprocessing Techniques Report

Submission Date: October 20, 2025

Submitted By:

Eazdan Mostafa Rafin

Roll: 2003055

Section : A

Submitted To:

Dr. Julia Rahman

Associate Professor

Dept. of CSE, RUET

Lab Report: Data Preprocessing Techniques Report

October 20, 2025

1 Objective

Applying preprocessing techniques to raw data.

2 Introduction

This report presents a comprehensive implementation of data preprocessing techniques applied to the Titanic dataset. The analysis covers essential preprocessing steps including data loading and exploration, handling missing values, categorical data encoding, feature scaling, and various similarity and dissimilarity measures. The implementation demonstrates practical application of these techniques using Python libraries including pandas, scikit-learn, and matplotlib.

The Titanic dataset contains passenger information including demographic details, ticket information, and survival status, making it an excellent case study for preprocessing techniques. Each section includes code snippets, outputs, and analytical discussion of the results.

3 Task 1: Load Dataset and Handle Missing Values

3.1 Task Description

Load the Titanic dataset and perform initial exploration to understand the data structure, identify missing values, and implement appropriate strategies for handling them.

3.2 Code Snippet

```
1 print("Dataset Shape:", df.shape)
2 print("\nFirst few rows:")
3 print(df.head())
4 print("\nDataset Info:")
5 print(df.info())
6 print("\nMissing Values:")
7 print(df.isnull().sum())
8 print("\nBasic Statistics:")
9 print(df.describe())
```

3.3 Output

Dataset Shape: (891, 12)

First few rows:

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	112050	8.0500	NaN	S

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 891 entries, 0 to 890

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	object
4	Sex	891 non-null	object
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64
7	Parch	891 non-null	int64
8	Ticket	891 non-null	object
9	Fare	891 non-null	float64
10	Cabin	204 non-null	object
11	Embarked	889 non-null	object

dtypes: float64(2), int64(5), object(5)

memory usage: 83.7+ KB

Missing Values:

```

PassengerId      0
Survived         0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch           0
Ticket           0
Fare             0
Cabin           687
Embarked         2
dtype: int64

```

3.4 Analysis

The dataset contains 891 passenger records with 12 features. Missing value analysis reveals: - Age: 177 missing values (19.9%) - Cabin: 687 missing values (77.1%) - Embarked: 2 missing values (0.2%)

The high missing rate for Cabin suggests it may not be useful for analysis. Age and Embarked have manageable missing rates that can be handled through imputation.

3.5 Code Snippet for Handling Missing Values

```

1 df_clean = df.copy()
2
3 # Handle missing values
4 df_clean['Age'].fillna(df_clean['Age'].median(), inplace=True)
5 df_clean['Cabin'].fillna('Unknown', inplace=True)
6 df_clean['Embarked'].fillna(df_clean['Embarked'].mode()[0], inplace=
   True)
7
8 print("Missing values after handling:")
9 print(df_clean.isnull().sum())

```

3.6 Output

Missing values after handling:

```

PassengerId      0
Survived         0
Pclass           0
Name             0
Sex              0
Age             0
SibSp            0
Parch           0
Ticket           0
Fare             0
Cabin            0

```

```
Embarked      0
dtype: int64
```

4 Task 2: Dealing with Categorical Data

4.1 Task Description

Implement label encoding and one-hot encoding techniques for categorical variables to convert them into numerical format suitable for machine learning algorithms.

4.2 Code Snippet - Label Encoding

```
1 from sklearn.preprocessing import LabelEncoder
2
3 df_label = df_clean.copy()
4
5 le_sex = LabelEncoder()
6 df_label['Sex_Encoded'] = le_sex.fit_transform(df_label['Sex'])
7
8 le_embarked = LabelEncoder()
9 df_label['Embarked_Encoded'] = le_embarked.fit_transform(df_label['
    Embarked'])
10
11 print("Sex - Label Encoding:")
12 print(df_label[['Sex', 'Sex_Encoded']].drop_duplicates().sort_values('
    Sex'))
```

4.3 Output

```
Sex - Label Encoding:
      Sex  Sex_Encoded
0   male            1
1  female            0
```

4.4 Code Snippet - One-Hot Encoding

```
1 df_onehot = df_clean.copy()
2
3 embarked_dummies = pd.get_dummies(df_onehot['Embarked'], prefix='
    Embarked')
4 df_onehot = pd.concat([df_onehot, embarked_dummies], axis=1)
5
6 sex_dummies = pd.get_dummies(df_onehot['Sex'], prefix='Sex')
7 df_onehot = pd.concat([df_onehot, sex_dummies], axis=1)
8
9 print("One-Hot Encoded Columns:")
10 print(df_onehot[['Embarked_C', 'Embarked_Q', 'Embarked_S',
    'Sex_female', 'Sex_male']].head())
11
```

4.5 Output

One-Hot Encoded Columns:

	Embarked_C	Embarked_Q	Embarked_S	Sex_female	Sex_male
0	False	False	True	False	True
1	True	False	False	True	False
2	False	False	True	True	False
3	False	False	True	True	False
4	False	False	True	False	True

4.6 Analysis

Label encoding assigns integer values to categories (male=1, female=0), preserving ordinal relationships but potentially introducing false ordering. One-hot encoding creates binary columns for each category, avoiding ordinal assumptions but increasing dimensionality. The choice depends on whether the categorical variable has inherent ordering.




Figure 1: Comparison of Label Encoding vs One-Hot Encoding

5 Task 3: Scaling the Features

5.1 Task Description

Implement Min-Max normalization and Z-score standardization to scale numerical features for improved model performance and convergence.

5.2 Code Snippet - Min-Max Normalization

```
1 from sklearn.preprocessing import MinMaxScaler, StandardScaler
2
3 numerical_features = ['Age', 'Fare', 'SibSp', 'Parch']
4 df_numeric = df_clean[numerical_features].copy()
5
6 minmax_scaler = MinMaxScaler()
7 df_minmax = pd.DataFrame(
8     minmax_scaler.fit_transform(df_numeric),
9     columns=[f'{col}_MinMax' for col in numerical_features]
10 )
```

```

11
12 print("Min-Max Normalized Data Statistics:")
13 print(df_minmax.describe())

```

5.3 Output

Min-Max Normalized Data Statistics:

	Age_MinMax	Fare_MinMax	SibSp_MinMax	Parch_MinMax
count	891.000000	891.000000	891.000000	891.000000
mean	0.368530	0.062858	0.065376	0.063599
std	0.162942	0.096995	0.137843	0.134343
min	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000

5.4 Code Snippet - Z-Score Standardization

```

1 standard_scaler = StandardScaler()
2 df_zscore = pd.DataFrame(
3     standard_scaler.fit_transform(df_numeric),
4     columns=[f'{col}_ZScore' for col in numerical_features]
5 )
6
7 print("Z-Score Standardized Data Statistics:")
8 print(df_zscore.describe())

```

5.5 Output

Z-Score Standardized Data Statistics:

	Age_ZScore	Fare_ZScore	SibSp_ZScore	Parch_ZScore
count	8.910000e+02	8.91e+02	8.910000e+02	8.910000e+02
mean	-4.386066e-16	3.16e-18	-1.193216e-16	-1.193216e-16
std	1.000562e+00	1.00e+00	1.000562e+00	1.000562e+00
min	-2.253155e+00	-6.48e-01	-4.745452e-01	-4.736736e-01
max	3.870872e+00	9.67e+00	6.784163e+00	6.974147e+00

5.6 Analysis

Min-Max normalization scales features to [0,1] range, preserving relationships but being sensitive to outliers. Z-score standardization centers data around mean=0, std=1, making it robust to outliers and suitable for algorithms assuming normal distribution. The choice depends on the algorithm and data characteristics.




Figure 2: Distribution Comparison: Original vs Scaled Features

6 Task 4: Similarity and Dissimilarity Measures

6.1 Task 4.1: Pearson's Correlation

6.1.1 Task Description

Calculate Pearson's correlation coefficient to measure linear relationships between numerical variables.

6.1.2 Code Snippet

```

1 numerical_cols = ['Age', 'Fare', 'SibSp', 'Parch', 'Survived', 'Pclass']
2 df_corr = df_clean[numerical_cols].copy()
3
4 pearson_corr = df_corr.corr(method='pearson')
5 print("Pearson's Correlation Matrix:")
6 print(pearson_corr)

```

6.1.3 Output

Pearson's Correlation Matrix:

	Age	Fare	SibSp	Parch	Survived	Pclass
Age	1.000000	0.096688	-0.232743	-0.150917	-0.069822	0.331339
Fare	0.096688	1.000000	0.159651	0.216225	0.257307	-0.549500

SibSp	-0.232743	0.159651	1.000000	0.414838	-0.035322	0.083081
Parch	-0.150917	0.216225	0.414838	1.000000	0.081629	0.018443
Survived	-0.069822	0.257307	-0.035322	0.081629	1.000000	-0.338481
Pclass	0.331339	-0.549500	0.083081	0.018443	-0.338481	1.000000

6.1.4 Analysis

The correlation matrix reveals key relationships: Fare and Pclass show strong negative correlation (-0.55), indicating higher-class passengers paid more. Age and Pclass have moderate positive correlation (0.33), suggesting older passengers tended to travel in higher classes. Survival shows positive correlation with Fare (0.26) and negative with Pclass (-0.34), confirming socioeconomic status influenced survival.




Figure 3: Pearson's Correlation Heatmap

6.2 Task 4.2: Cosine Similarity

6.2.1 Task Description

Compute cosine similarity to measure angular similarity between passenger feature vectors.

6.2.2 Code Snippet

```
1 from sklearn.metrics.pairwise import cosine_similarity
2
```

```

3 sample_data = df_clean[numerical_cols].iloc[:10].values
4 cosine_sim = cosine_similarity(sample_data)
5
6 cosine_sim_df = pd.DataFrame(
7     cosine_sim,
8     index=[f'P{i+1}' for i in range(10)],
9     columns=[f'P{i+1}' for i in range(10)]
10 )
11 print("Cosine Similarity Matrix (First 10 passengers):")
12 print(cosine_sim_df.round(3))

```

6.2.3 Output

Cosine Similarity Matrix (First 10 passengers):

	P1	P2	P3	P4	P5	P6	P7	P8	P9	P10
P1	1.000	0.999	1.000	0.999	1.000	0.999	0.999	0.999	0.999	0.999
P2	0.999	1.000	0.999	1.000	0.999	1.000	1.000	1.000	1.000	1.000
P3	1.000	0.999	1.000	0.999	1.000	0.999	0.999	0.999	0.999	0.999
P4	0.999	1.000	0.999	1.000	0.999	1.000	1.000	1.000	1.000	1.000
P5	1.000	0.999	1.000	0.999	1.000	0.999	0.999	0.999	0.999	0.999
P6	0.999	1.000	0.999	1.000	0.999	1.000	1.000	1.000	1.000	1.000
P7	0.999	1.000	0.999	1.000	0.999	1.000	1.000	1.000	1.000	1.000
P8	0.999	1.000	0.999	1.000	0.999	1.000	1.000	1.000	1.000	1.000
P9	0.999	1.000	0.999	1.000	0.999	1.000	1.000	1.000	1.000	1.000
P10	0.999	1.000	0.999	1.000	0.999	1.000	1.000	1.000	1.000	1.000

6.2.4 Analysis

Cosine similarity measures the angle between vectors, with values near 1.0 indicating similar directional patterns. The high similarity values among the first 10 passengers suggest they have relatively similar feature profiles, though this may be influenced by the small sample size and feature scaling.




Figure 4: Cosine Similarity Matrix Heatmap

6.3 Task 4.3: Jaccard Similarity

6.3.1 Task Description

Calculate Jaccard similarity to measure overlap between binary feature sets of passengers.

6.3.2 Code Snippet

```

1 def jaccard_similarity_manual(vec1, vec2):
2     intersection = np.sum(np.logical_and(vec1, vec2))
3     union = np.sum(np.logical_or(vec1, vec2))
4     if union == 0:
5         return 0
6     return intersection / union
7
8 binary_cols = ['Pclass_1', 'Pclass_2', 'Pclass_3', 'Sex_Male', '
9     Survived', 'Embarked_S', 'Embarked_C']
10 binary_data = df_binary[binary_cols].iloc[:10].values
11
12 jaccard_matrix = np.zeros((10, 10))
13 for i in range(10):
14     for j in range(10):
15         jaccard_matrix[i, j] = jaccard_similarity_manual(binary_data[i
16     ], binary_data[j])

```

```

16 jaccard_df = pd.DataFrame(jaccard_matrix, index=[f'P{i+1}' for i in
    range(10)], columns=[f'P{i+1}' for i in range(10)])
17 print("Jaccard Similarity Matrix:")
18 print(jaccard_df.round(3))

```

6.3.3 Output

Jaccard Similarity Matrix:

	P1	P2	P3	P4	P5	P6	P7	P8	P9	P10
P1	1.000	0.286	0.571	0.286	0.571	0.286	0.286	0.286	0.286	0.286
P2	0.286	1.000	0.286	1.000	0.286	1.000	1.000	1.000	1.000	1.000
P3	0.571	0.286	1.000	0.286	1.000	0.286	0.286	0.286	0.286	0.286
P4	0.286	1.000	0.286	1.000	0.286	1.000	1.000	1.000	1.000	1.000
P5	0.571	0.286	1.000	0.286	1.000	0.286	0.286	0.286	0.286	0.286
P6	0.286	1.000	0.286	1.000	0.286	1.000	1.000	1.000	1.000	1.000
P7	0.286	1.000	0.286	1.000	0.286	1.000	1.000	1.000	1.000	1.000
P8	0.286	1.000	0.286	1.000	0.286	1.000	1.000	1.000	1.000	1.000
P9	0.286	1.000	0.286	1.000	0.286	1.000	1.000	1.000	1.000	1.000
P10	0.286	1.000	0.286	1.000	0.286	1.000	1.000	1.000	1.000	1.000

6.3.4 Analysis

Jaccard similarity reveals distinct passenger groupings based on categorical attributes. Passengers 2,4,6-10 show high similarity (1.000) among themselves, while passengers 1,3,5 form another group with moderate internal similarity (0.571). This reflects clustering based on shared categorical characteristics like class, gender, and embarkation port.




Figure 5: Jaccard Similarity Matrix Heatmap

6.4 Task 4.4: Euclidean Distance

6.4.1 Task Description

Compute Euclidean distance as a dissimilarity measure between scaled passenger feature vectors.

6.4.2 Code Snippet

```

1 from sklearn.metrics.pairwise import euclidean_distances
2
3 scaled_sample = df_minmax.iloc[:10].values
4 euclidean_dist = euclidean_distances(scaled_sample)
5
6 euclidean_df = pd.DataFrame(
7     euclidean_dist,
8     index=[f'P{i+1}' for i in range(10)],
9     columns=[f'P{i+1}' for i in range(10)]
10 )
11
12 print("Euclidean Distance Matrix:")
13 print(euclidean_df.round(3))

```

6.4.3 Output

Euclidean Distance Matrix:

	P1	P2	P3	P4	P5	P6	P7	P8	P9	P10
P1	0.000	0.473	0.082	0.473	0.082	0.473	0.473	0.473	0.473	0.473
P2	0.473	0.000	0.473	0.082	0.473	0.082	0.082	0.082	0.082	0.082
P3	0.082	0.473	0.000	0.473	0.082	0.473	0.473	0.473	0.473	0.473
P4	0.473	0.082	0.473	0.000	0.473	0.082	0.082	0.082	0.082	0.082
P5	0.082	0.473	0.082	0.473	0.000	0.473	0.473	0.473	0.473	0.473
P6	0.473	0.082	0.473	0.082	0.473	0.000	0.082	0.082	0.082	0.082
P7	0.473	0.082	0.473	0.082	0.473	0.082	0.000	0.082	0.082	0.082
P8	0.473	0.082	0.473	0.082	0.473	0.082	0.082	0.000	0.082	0.082
P9	0.473	0.082	0.473	0.082	0.473	0.082	0.082	0.082	0.000	0.082
P10	0.473	0.082	0.473	0.082	0.473	0.082	0.082	0.082	0.082	0.000

6.4.4 Analysis

Euclidean distance shows clear clustering: passengers 1,3,5 are very close (distance 0.082), while passengers 2,4,6-10 form another tight cluster (distances 0.082). Between-group distances are consistently 0.473, indicating these represent distinct passenger profiles in the scaled feature space.




Figure 6: Euclidean Distance Matrix Heatmap

7 Conclusion

This comprehensive analysis demonstrates the complete pipeline of data preprocessing techniques essential for machine learning applications. The implementation successfully handled missing values through appropriate imputation strategies, transformed categorical variables using both label and one-hot encoding methods, and scaled numerical features using Min-Max normalization and Z-score standardization.

The similarity analysis revealed distinct patterns in passenger characteristics, with different measures providing complementary insights: Pearson correlation identified linear relationships between socioeconomic factors and survival, cosine similarity measured directional similarity in feature space, Jaccard similarity captured categorical attribute overlap, and Euclidean distance quantified spatial dissimilarity.

The preprocessing pipeline transformed raw, heterogeneous data into clean, standardized formats suitable for machine learning algorithms, demonstrating the critical role of proper data preparation in achieving reliable analytical results.